

# Report for Music Interpolation

Krishna	1001781
Gao Yunyi	1002871
Pengfei	1002949

## Abstract


In this report, we aim to interpolate between 2 music excerpts using the extracted feature of mel frequency. Our main tasks are to use autoencoder to get latent vectors of given music excerpts, interpolate between two latent vectors, decode the interpolated latent vector to mel spectrogram and reconstruct to audio. Our experiment is conducted in stages including dataset collection, data preprocessing, experimentation, evaluation and conclusion, which will be presented in the following report.

## Dataset and Collection


We use the open-source audio dataset called Free Music Archive (FMA). It is an interactive library of high-quality, legal audio downloads directed by WFMU launched in 2009 (<https://github.com/mdeff/fma>). It contains MP3 files of 77,643 songs, which are 30 seconds each covering across 68 genres and sampled at 44,100Hz, 128kb/s. It contains not only the audio but also precalculated feature of each song as below.

Low-level Musical Features		High-level Musical Features	Other Metadata
Feature	Dimensions	Acousticness Danceability Energy Instrumentalness Liveness Speechiness Tempo Valence	Genre Album Artist
Chroma_Cens	83		
Chroma_CQT	83		
Chroma_STFT	83		
MFCC	139		
RMSE	6		
Spectral Bandwidth	6		
Spectral Centroid	6		
Spectral Contrast	48		
Spectral Rolloff	6		
Tonex	41		
ZCR	6		

The following is our graphical analysis of the genre distribution over 16 top genres in total. Number is the amount of music tracks under the stated genre.

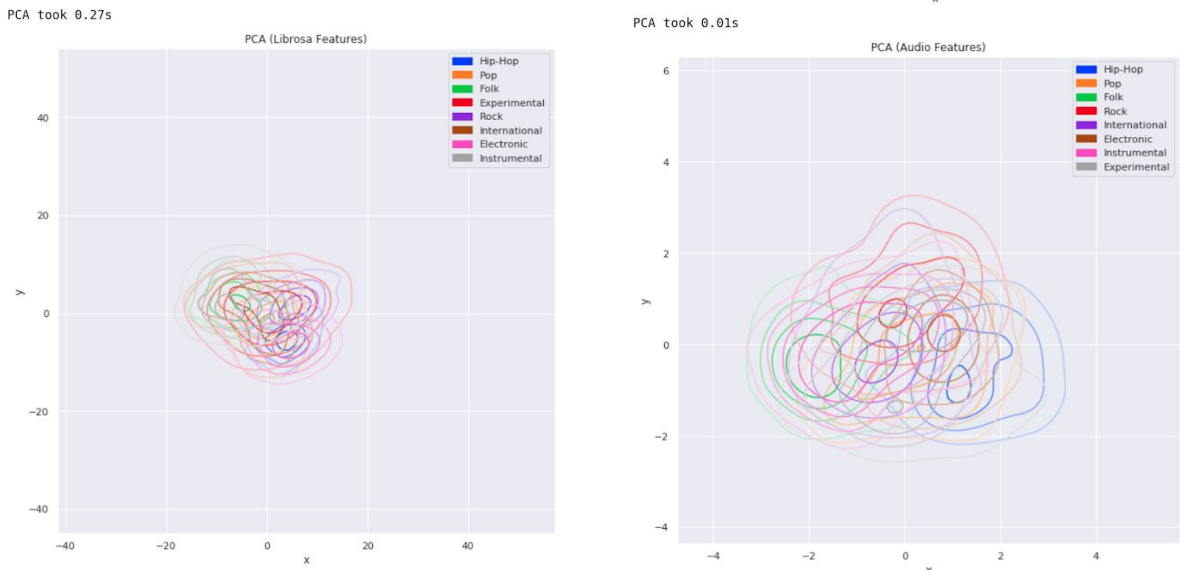


Due to the unbalanced dataset with respect to genre, FMA has already defined a balanced data subset over 8 top genres in its metadata. We have filtered out the balanced subset and below is our graphical visualization of the balanced subset which is used for our future training.

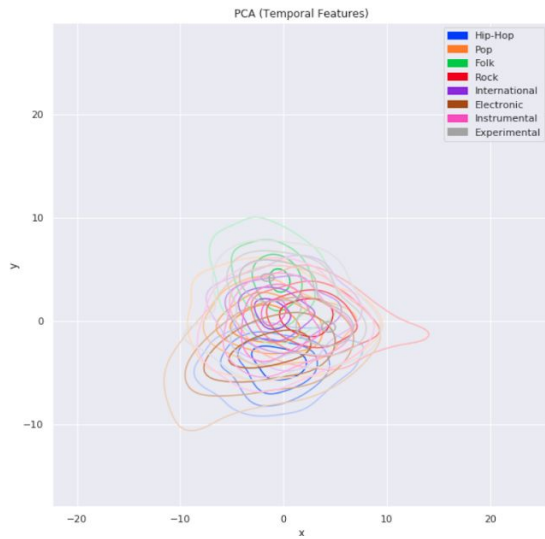


### Visualization of low-level musical features

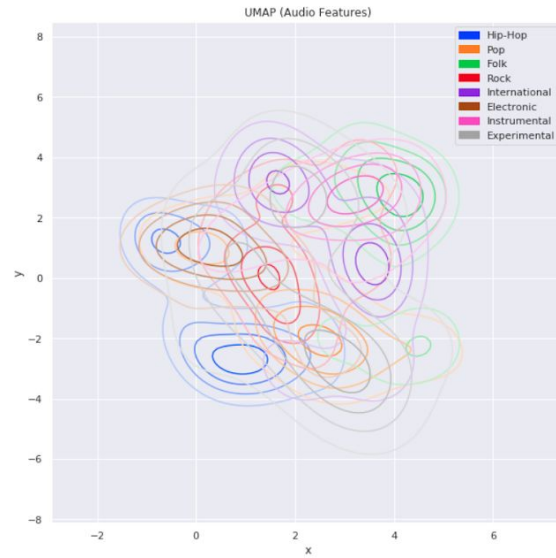
We visualize the balanced dataset distribution over 3 subset of low-level features: librosa feature, temporal feature and audio feature. Three dimensionality reduction methodologies are used including: Tsne, PCA and UMAP.




PCA took 0.02s



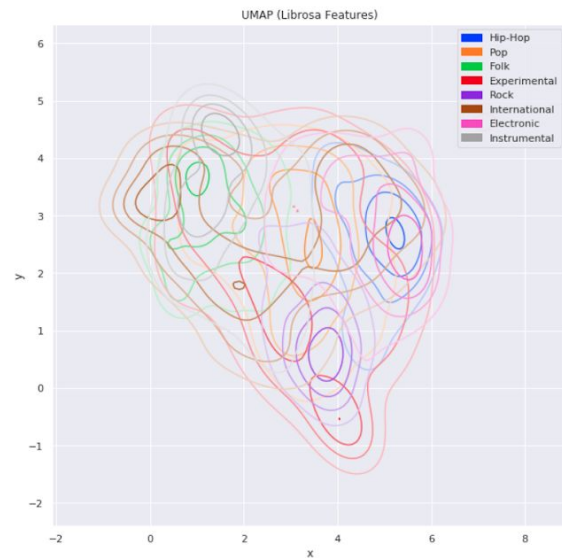
UMAP took 5.36s



UMAP took 3.26s



UMAP took 101.00s




## Visualization of high-level musical features

The combination of high level musical features including Acousticness, Danceability, Energy, Instrumentalness, Liveness, Speechiness, Tempo, Valence are visualized with respect to genre using 2 methods: Tsne and PCA, which the latter gives clearer separation between genres.

### T-sne on genre distributions

```
plt.scatter(X_tsne[:, 0], X_tsne[:, 1], c=[y2c[o] for o in y.values], s=0.5)
```


```
<matplotlib.collections.PathCollection at 0x7fedf4231810>
```



### PCA on genre distributions

```
plt.scatter(X_PCA[:, 0], X_PCA[:, 1], c=[y2c[o] for o in y.values], s=1)
```


```
<matplotlib.collections.PathCollection at 0x7fedf444b3d0>
```




## Data Pre-processing

In order to interpolate between audio samples, we need a generative model that's able to capture both the low-level and high-level features of audio clips. Earlier works such as WaveNet aim to do this by directly modelling the individual samples in audio clips. Since audio in its raw form requires hundreds of thousands of samples to represent just a few seconds, and due to the inherently serial nature of audio, modelling it in such a manner means that the learning algorithm would need to learn the relationships between samples that are hundreds of thousands of time-steps apart.

In order to enable our model(s) to capture higher-level information about audio clips, we convert the time-domain audio signal into a time-frequency domain spectrogram. We do this by sequentially applying an FFT over overlapping windows of the time-domain audio signal to obtain what's known as a Short-Time Fourier Transform (STFT).



To better represent the human perception of pitch, we rescale the frequency axis to use the Mel scale. In addition, we take the logarithm of each value in the spectrogram to more closely match how humans perceive loudness.



Since our goal is to be able to generate audio samples (and not just their spectrogram representations), we also need to be able to reconstruct audio clips from Mel spectrograms. Unfortunately, the reconstruction is unavoidably lossy because in deriving a Mel spectrogram we discard some temporal information (FFT Phase) and some frequency information (in converting

to the Mel scale).

To convert a Mel spectrogram to audio, we must first recover the removed frequency information by converting the Mel scale to a linear frequency scale. Thereafter, we reconstruct the phase information and apply the inverse FFT function to recover the audio. The Griffin-Lim algorithm is one method of estimating the FFT phase array based on the FFT magnitude array.

The Small subset of the FMA Dataset contains around ~8000 audio clips trimmed to approximately 30 seconds each. After removing ~10 clips that were malformed or unusable, we converted each one to its Mel spectrogram representation. We then split each spectrogram into subsections of 640 frames each, which with our FFT parameters corresponds to 5 seconds of audio. Splitting the spectrograms in this manner and discarding any subsections shorter than the specified 640 frames gives us a clean, consistent dataset. The genre-balanced nature of the small subset of the FMA Dataset we use means that we do not have to perform any rebalancing.

We use `librosa` along with `numpy` to perform the majority of our data pre-processing. To ensure portability and to keep the disk-footprint of the generated spectrograms small, we store the generated spectrograms as OpenEXR images. OpenEXR is a lossless open image standard that supports 32-bit floating-point values, and is therefore ideal for holding spectrogram data.

When converting an audio clip to a Mel spectrogram, we note a tradeoff between reconstruction quality and reconstruction time. Using `librosa`, the spectrograms resulting from the parameters we chose take about 18 seconds to reconstruct (per second of audio) on a high-end consumer CPU.

The parameters we use for spectrogram generation are as follows:

Parameter	Value
Audio Sampling Rate	32768 Hz
FFT Window Size	256
STFT Hop Length	4096
Mel Frequency Bins	512

## Problem and Algorithm / Model

Inspired by MusicVAE (<https://magenta.tensorflow.org/music-vae>), a machine learning model that let us create palettes for blending and exploring musical scores using MIDI files, we intend to make a model that interpolate between two music excerpts using mel spectrogram, which describes richer content of the audio information.

MusicVAE used one latent space model called autoencoder (AE) which is capable of learning the fundamental characteristics of a training dataset. An autoencoder builds a latent space of a dataset by learning to compress (encode) each example into a latent vector, and then reproduce (decode) the same example from that latent vector. MusicVAE conducts reconstruction, sampling, and smooth interpolation based on this architecture.



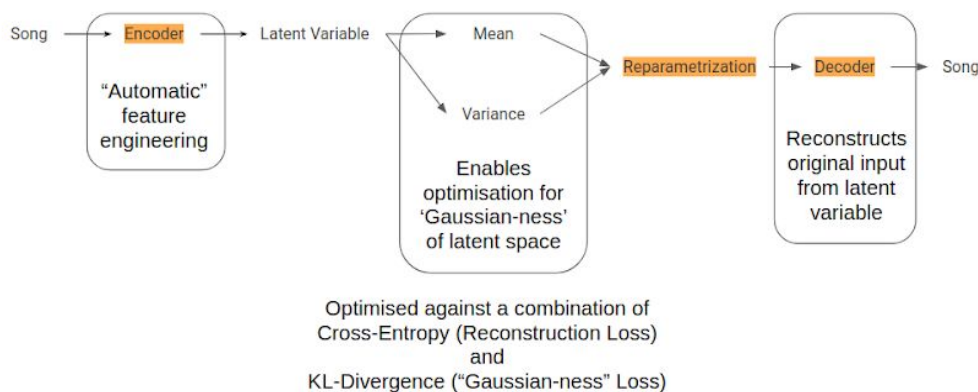
## Our Approach

Our tasks include:


- 1) unconditioned audio generation
- 2) audio interpolation

### 1) Variational Autoencoder (VAE) for latent space interpolation

In order to interpolate music, we need to explore on the variations of our original dataset. The standard autoencoder generates encodings specifically useful for reconstructing its own input using CNN. The loss function is usually either the mean-squared error or cross-entropy between the output and the input, known as the reconstruction loss, which penalizes the network for creating outputs different from the input. While VAE, not making its encoder not output an encoding vector of size  $n$  like standard autoencoder, rather, outputting two vectors of size  $n$ : a vector of means,  $\mu$ , and another vector of standard deviations,  $\sigma$ . Hence, decoder learns that not only is a single point in latent space referring to a sample of that class, but all nearby points refer to the same as well. This allows the decoder to not just decode single, specific encodings in the latent space, but ones that slightly vary too.



### 2) Latent Space Interpolation




The interpolation of two music excerpts is done by linear algebra of the latent space learned by task(1) of music excerpt A and B and do a linear combination of encoded A and B latent vectors, the constant of "fraction" in the above equation is a scalar determining the magnitude and direction of the interpolation.

The following records our different experiments on building and improving the above 2 tasks.

## Experiments

### 1) Baseline model

In our first attempt, we used CNN as encoder.






Model: "Encoder"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 255, 319, 32)	320
conv2d_1 (Conv2D)	(None, 127, 159, 64)	18496
conv2d_2 (Conv2D)	(None, 63, 79, 64)	36928
conv2d_3 (Conv2D)	(None, 31, 39, 64)	36928
flatten (Flatten)	(None, 77376)	0
dense (Dense)	(None, 64)	4952128
Total params: 5,044,800		
Trainable params: 5,044,800		
Non-trainable params: 0		


Model: "Decoder"

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 81920)	2703360
reshape (Reshape)	(None, 32, 40, 64)	0
conv2d_transpose (Conv2DTran	(None, 64, 80, 64)	36928
conv2d_transpose_1 (Conv2DTr	(None, 128, 160, 64)	36928
conv2d_transpose_2 (Conv2DTr	(None, 256, 320, 64)	36928
conv2d_transpose_3 (Conv2DTr	(None, 512, 640, 32)	18464
conv2d_transpose_4 (Conv2DTr	(None, 512, 640, 1)	289
Total params: 2,832,897		
Trainable params: 2,832,897		
Non-trainable params: 0		

This gives results with one sample of reconstructed spectrogram from its latent vector as below:

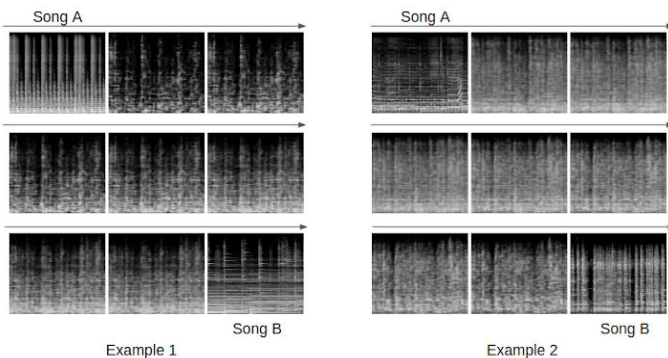


Original spectrogram



Reconstructed spectrogram

Below are two samples of interpolated spectrogram from 2 spectrogram reconstructed from baseline model.



In the baseline model, the reconstruction result is not very similar to the original result.

Through tuning the parameters, we have found out that we are not able to reconstruct on a single model, after debugging, we found out the loss function we have applied before -- `tf.nn.sigmoid_cross_entropy_with_logits` first maps a logit in sigmoid, which is (0, 1) and then do minimization on cross entropy loss. This works well on 0 or 1 digits classification, but however, it does not work well on output values that range bigger than 1. So in the end we have applied the mean squared entropy loss to minimize the loss. After which we are able to overfit on one sample.

## 2)Conv2D

### 3) Conv2D inception

The baseline model used CNN. Audio spectrograms are inherently sequential, which means pattern in the lower-left corner of a spectrogram does not mean the same thing as an identical pattern in the top-right corner. CNNs are translation-invariant, whereas spectrograms are not. A better model would take into account the fact that patterns in different locations on the spectrogram have different meanings (pitch, time, etc.) Hence in the following attempts, we used Cov1D to mitigate this problem.

### 3) Conv1D

In Conv1d, our original input shape is (batch\_size, frequency\_dimension, time\_dimension), we have reshaped the

## Evaluation Methodology

Our task of interpolation of two music excerpts can be viewed in 2 main tasks:

(1) unconditioned audio generation (2) audio interpolation

Evaluation Metrics for unconditioned audio generation contains the following questions which are asked in the survey:

Is the output fluent on transitions?

Does the output sound like before reconstructing?

Does the music sound good?

Evaluation Metrics for audio interpolation contains the following questions which are asked in the survey:

Does the output sound like a combination of the two inputs?

Are the features preserved after interpolating?

Are the genre preserved after interpolating?

Genre conservation is tested using genre classifier trained over FMA dataset. Below are our different attempts of genre classifier.

### Genre Classifier

#### Machine Learning approach:

(different parameters are tried given same architecture, only the best result shown below)

#### 1)Decision Tree

leaf\_ls=[1,3,5,7,10,25]

min leaves=10				
	precision	recall	f1-score	support
0	0.31	0.25	0.28	297
1	0.33	0.29	0.31	481
2	0.20	0.36	0.26	146
3	0.23	0.18	0.20	277
4	0.17	0.36	0.23	134
5	0.32	0.19	0.24	312
6	0.17	0.14	0.15	304
7	0.34	0.50	0.41	129
accuracy			0.26	2000
macro avg	0.26	0.28	0.26	2000
weighted avg	0.27	0.26	0.25	2000

#### 2)Logistic Regression

C\_param\_range = [0.001,0.01,0.1,1,10,100]

Regularization parameter(C)=0.1				
	precision	recall	f1-score	support
0	0.41	0.34	0.37	297
1	0.32	0.17	0.22	481
2	0.37	0.47	0.41	146
3	0.38	0.42	0.40	277
4	0.23	0.60	0.33	134
5	0.55	0.40	0.47	312
6	0.22	0.20	0.21	304
7	0.43	0.64	0.51	129
accuracy			0.35	2000
macro avg	0.36	0.40	0.37	2000
weighted avg	0.37	0.35	0.35	2000

#### 3)SVM

C\_param\_range = [0.001,0.01,0.1,1,10,100,150,200]

Regularization parameter(C)=100				
	precision	recall	f1-score	support
0	0.53	0.60	0.56	328
1	0.54	0.44	0.48	380
2	0.67	0.67	0.67	431
3	0.63	0.66	0.65	316
4	0.61	0.65	0.63	456
5	0.72	0.58	0.64	340
6	0.48	0.52	0.50	426
7	0.61	0.62	0.62	323
accuracy			0.59	3000
macro avg	0.60	0.59	0.59	3000
weighted avg	0.60	0.59	0.59	3000

#### 4)KNN

n\_neighbours= [5,10,20,40]

n_neighbors=40				
	precision	recall	f1-score	support
0	0.40	0.26	0.32	297
1	0.25	0.13	0.18	401
2	0.39	0.49	0.43	146
3	0.31	0.49	0.38	277
4	0.22	0.46	0.30	134
5	0.51	0.46	0.48	312
6	0.23	0.15	0.18	304
7	0.40	0.64	0.49	129
accuracy			0.34	2000
macro avg	0.34	0.39	0.35	2000
weighted avg	0.34	0.34	0.32	2000

#### 5)Guassian Naive Bayes

	precision	recall	f1-score	support
0	0.24	0.05	0.09	297
1	0.36	0.04	0.07	401
2	0.15	0.45	0.23	146
3	0.31	0.50	0.39	277
4	0.28	0.67	0.39	134
5	0.31	0.06	0.10	312
6	0.19	0.10	0.13	304
7	0.24	0.91	0.38	129
accuracy			0.24	2000
macro avg	0.26	0.35	0.22	2000
weighted avg	0.27	0.24	0.18	2000

#### 6)GMM

without kmeans mean init

	precision	recall	f1-score	support
0	0.00	0.00	0.00	297
1	0.25	0.02	0.04	401
2	0.00	0.00	0.00	146
3	0.00	0.00	0.00	277
4	0.00	0.00	0.00	134
5	0.00	0.00	0.00	312
6	0.15	0.99	0.27	304
7	0.00	0.00	0.00	129
accuracy			0.15	2000
macro avg	0.05	0.13	0.04	2000
weighted avg	0.07	0.15	0.05	2000

with kmeans init (means initialized by one sampler by genre)

	precision	recall	f1-score	support
0	0.15	1.00	0.26	297
1	0.00	0.00	0.00	401
2	0.00	0.00	0.00	146
3	0.00	0.00	0.00	277
4	0.00	0.00	0.00	134
5	0.00	0.00	0.00	312
6	0.00	0.00	0.00	304
7	0.00	0.00	0.00	129
accuracy			0.15	2000
macro avg	0.02	0.12	0.03	2000
weighted avg	0.02	0.15	0.04	2000

### Deep Learning approach:

(Train on 35000 samples, validate on 8320 sample over 43000 samples in balanced dataset)

#### 1)CNN\_1:

Model: "sequential\_6"

Layer (type)	Output Shape	Param #
conv2d_18 (Conv2D)	multiple	160
conv2d_19 (Conv2D)	multiple	2320
conv2d_20 (Conv2D)	multiple	2320
max_pooling2d_6 (MaxPooling2D)	multiple	0
dropout_6 (Dropout)	multiple	0
flatten_6 (Flatten)	multiple	0
dense_18 (Dense)	multiple	167772288
dense_19 (Dense)	multiple	8256
dense_20 (Dense)	multiple	520
Total params: 167,785,864		
Trainable params: 167,785,864		
Non-trainable params: 0		

## Results:

Epoch Number	Train Accuracy	Validation Accuracy
1	0.4092	0.4099
2	0.6301	0.4128
3	0.7038	0.4012
4	0.8396	0.3960
5	0.9348	0.3952

Over the epochs, training accuracy increases quickly but validation accuracy remains at around 0.4 and drops after 2nd epochs. The model will be del is overfitting probably due to overly deep architecture. Fewer layers used in next attempt.

## 2)CNN\_2:

Layer (type)	Output Shape	Param #
conv2d_38 (Conv2D)	multiple	303
conv2d_39 (Conv2D)	multiple	4515
conv2d_40 (Conv2D)	multiple	97565
flatten_12 (Flatten)	multiple	0
dense_28 (Dense)	multiple	2017608
Total params: 2,119,991		
Trainable params: 2,119,991		
Non-trainable params: 0		

## Results:

Epoch Number	Train Accuracy	Validation Accuracy
1	0.2664	0.3132

2	0.3903	0.3463
3	0.4623	0.4303
4	0.5096	0.4547
5	0.5533	0.4480
6	0.5969	0.4561
7	0.6573	0.4301
8	0.7309	0.4104
9	0.8925	0.3919
10	0.9712	0.3849

Over the epochs, training accuracy increases quickly over the epochs but validation accuracy remains at around 0.4 and drops after 6th epochs. Con2D may not work for time-series features and we will attempt on RNN (recurrent neural network), which may fit better for music features.

### 3)RNN:

---

Model: "sequential_1"		
Layer (type)	Output Shape	Param #
=====		
lstm_2 (LSTM)	(None, 512, 128)	393728
lstm_3 (LSTM)	(None, 512, 128)	131584
dropout_1 (Dropout)	(None, 512, 128)	0
time_distributed_4 (TimeDist	(None, 512, 64)	8256
time_distributed_5 (TimeDist	(None, 512, 32)	2080
time_distributed_6 (TimeDist	(None, 512, 16)	528
time_distributed_7 (TimeDist	(None, 512, 8)	136
flatten_1 (Flatten)	(None, 4096)	0
dense_9 (Dense)	(None, 8)	32776
=====		
Total params: 569,088		
Trainable params: 569,088		
Non-trainable params: 0		

---

### Results:

Epoch Number	Train Accuracy	Validation Accuracy
1	0.3264	0.3681
2	0.4416	0.3927
3	0.4846	0.4028
4	0.5230	0.4106
5	0.5567	0.3972
6	0.5788	0.4038

7	0.6064	0.3950
8	0.6278	0.3679
9	0.6501	0.3791
10	0.6689	0.3668

Over the epochs, training accuracy increases over the epochs but validation accuracy remains at around 0.4 and drops after 4th epochs. Overfit issue also occurs.

Hence, given above attempts, SVM( with regularization parameter C equals to 100 gives the best validation accuracy)

## Results and Discussion